

Motion Detection with an HC-SR04 ultrasonic sensor: Issues, Statistics, and Thoughts

Gary Drykacz,

November, 2020

Introduction

The project goal was to explore using a programmable LED strip as a visual aid to guide a vehicle into a specifically defined area within a garage. Two sensor would be used, one pointed at the front of the car, and the second laterally on the side wall. The forward sensor would change the LED strip color based on distance to the garage back wall, and the led pattern would shift left, right or from edges to center in traveling waves, based on the side sensor.

We have a two car garage, with a relatively narrow "bay" for each car. If one car is parked too close to the edge of its bay, the driver of the other vehicle may not be able to easily get in or out of their car. In addition, it would be nice, but less critical, to have a well-defined consistent front stopping area. Of course, it was not practical to have the LED lights on continuously, so there are two criteria to analyze, motion and distance.

To be honest, long familiarity with parking the car in the garage worked fine, and an "upgrade" to a cheap, commercial, laser position device that projected a laser dot on the car also helped. The problem with the latter is that it showed too late to really position the car earlier in parking. Of course, there is the tried and true, cheap ball and string indicator, as well as other electronic methods to indicate final car position, but this more refined idea peaked my interest.

To jump to the project conclusion, the project as originally envisioned failed. The reasons for the failure where not the project concept, but a naive understanding of the ultrasonic echo response from the front bumper and grill of a modern vehicle, along with an unforeseen relationship of our driveway configuration with respect to the garage entryway. The latter problem was connected with misleading information on the HC-SR04 beam spread angle.

Even the vehicle with the least aerodynamic front was not capable of returning an echo with sufficient energy back to the sensor microphone to trigger a response, except at distances just over 2 meters.

Our driveway is "L" shaped with a tight turn into the garage. I thought the detection spread of the sensor would be sufficient to compensate for the angle in which the cars enter the garage. That assumption was overly optimistic. It was not until very late in the project I had access to laptop to monitor a vehicle entering the garage. A few tests quickly showed that the project in its original formulation was a non starter. Anyone anticipating a similar project, with curved entry path close to the garage entrance will likely have the same issues. For some vehicles and some driveway configurations, the original idea might work, but it is doubtful.

With a failure of the project, why even bother going into a very lengthy discussion of results? Four reasons:

1. Those wishing to do something similar, may realize the pitfalls to overcome.
2. Measuring motion, particular when the velocity may vary, is not the same as measuring distance, there are more considerations. Anyone needing to monitor motion and position in a project, may find the discussion here a useful starting point.
3. Sometimes failures can be just as useful as positive results.
4. Along the way on this project, I learned some interesting facts on how the sensor senses return echoes that might clear up some other information found online.

First, let's look at some general considerations of sensors that will affect this kind of project. Every sensor has an envelope of characteristics to negotiate for optimal use. For measuring distance and motion, we have several characteristics that need to be maneuvered through to find the best fit for a project:

- Lower and upper distance limits of detection,
- Error limits of individual measurements,
- Data repetition rate, i.e., intervals between readings versus total time of monitoring,
- Velocity difference between the sensor and target object,
- Characteristics of the target: size, shape, and texture.

In this project, we need to track two parameters. 1. The distance of the vehicle from the final resting position both frontally and laterally, and 2. Some factor that tells us if the car is moving or stopped. In addition, we want the display system activated only when a car is moving, not at rest, or not in the range (not in garage).

Tracking distance is straightforward. We know from the HC-S04 specifications that we can measure the distance to an object to about 0.5 cm, which is well below what we need for the current case.

Tracking motion, is not so easy to deal with. Motion measurements are derived from differences between successive values. More variables need to be negotiated: how fast we can take repetitive measurements, the velocity of the object, and the statistics of the sensor's echo time error limits, and total time available to make measurements. How to consider the interplay of these characteristics drives the majority of the discussion here. Monitoring has to be fast enough to measure the distance, and change the LED lights color and pattern.

We need to first look at the general mix of sensors that deal with motion and/or distance, as well as cost. Do you want to know if an object is just moving at all? If so, there are many choices for sensors: PIR; ultrasonic; laser; radar/microwaves. Second question: Do you also need to know the distance from the sensor? That narrows the sensor field to ultrasonic, laser, or multiple sensors, one for distance, one for motion. Last question: What range are we talking about? Up to a meter, way over a meter, like > 9 meters? For non precise measurements and measurements up to just over a meter, there are many cheap LIDAR/ToF (time of flight) modules to choose from. Beyond a meter, laser sensor prices take a leap. A great possibility, with excellent long range, is the, TF-Luna, or TFMMini variants, but we are talking about starting prices of \$25 and higher. Do you want to get away really cheap, and still measure long and short ranges? For distance only, the choice is an ultrasonic sensor, with the HC-SR04 being one of the most popular and well characterized. However, adding the need to measure motion, places constraints on an ultrasonic sensor.

Every sensor has drawbacks. For instance, LIDAR/ToF sensitivity will fall off rapidly with very dark objects, because they are based on near infrared light, which may or may not be absorbed by dark or black objects. Ultrasonic sensors will have issues with soft, highly porous materials, such as cloth. Both types of sensors will have problems with very rough textured surfaces, where diffuse reflectance may dominate over specular reflectance, and rounded or highly irregular shapes, which minimize any reflections.

All sensors have minimum and maximum range limits, We need to have some idea of the range limits for the project:

1. What is the maximum distance limit the sensor will start to monitor a velocity change? For ultrasonic sensors, the maximum is usually around 400 - 600 cm. Because of the speed of sound, this distance critically decides the lower limit for the ping interval, i.e., the time between successive pings. We throw values above this limit away. In principle, this range is just about right for the length of typical garage.

2. Range limits also depend on the time interval of the event, in this instance the time and distance to park a car in the garage. What is the anticipated velocity of a car entering the garage? For many garages, the distance from the garage entry to the back of the garage will be 18-22 ft, but maybe there is all sorts of stuff in the back of the garage. Use 18 ft, which is 5.49 m, as the total range. We also need to know how long it takes to park the car in the garage. We do a few timed measurements to park the car. Even manual counting: "1000, 2000, 3000, ..." is good enough in this case; even better, is recording the time with a stopwatch or smart phone. (As we will see later, using an HC-SR04 sensor, is generally not going to work for monitoring the time.) We find the time is between 8-11 s. Well, because the goal is not to see how fast we can park a car in the garage, let us say 11 s is a better time limit to use. The average velocity of the car is then around 50 cm/s. We have to be cautious about this average, a car's velocity will clearly not be constant through the 18 ft. However, for a first approximation of the issues, we will assume it is constant.

We have to worry about the relationship of the car's velocity, and number of pings we can do in the time frame we found we need to take measurements. The lower limit for motion detection may be bounded by the inherent error in sensor measurements for individual responses. If the difference between two successive measurements is below the error limit, we cannot be sure if motion has occurred. As a quick analysis, take the example of a HC-SR04 ultrasonic sensor. The quoted specification on HC-SR04 error is 0.5 cm. In terms of the speed of sound time difference, this is a round trip time from transmitting transducer, and back to the microphone transducer of 29 μ s. However, that is not the limiting factor for the ping interval time. The common ping interval is a thousand times higher, 33 ms. The reason is the speed of sound, ~ 0.344 m/ms and the 4-6 meter maximum range of an HC-SR04 sensor, e.g., 5.7 m, or 18.6 ft. Therefore, a guesstimate of the minimum velocity to just be able to detect movement outside the published error of the sensor is 0.5 cm/0.033s, or 15.2 cm/s (6 in/s). If the object velocity is less than 15.2 cm/s, we cannot be certain if the object has moved by comparing the difference between just two pings at a ping interval of 33 ms. Of course, we can increase the ping interval to unequivocally register a motion change, but that bumps up against another concern, the number of pings we can do within the total time frame for monitoring, which we found was 11 seconds. We can do 333 pings in a time limit of 11 seconds, which is a rate of 30.3 pings per second. so we do have some leeway to increase the ping interval. A large part of the work described here, will refine this quick analysis.

To determine the maximum velocity limit of the object, we have to decide the minimum pings that we can tolerate to adequately monitor the car's position. If we say that over the 5.5 meters in 11 seconds, we are willing to go as low as 55 pings, which is a measurement every 10 cm, then the upper velocity bound for monitoring the car is $(5.5\text{m/s} / 55 \text{ pings} \times 0.033 \text{ s})$ or 3.03 m/s (9.9 ft/s). Hopefully, that is quite a bit faster than most people enter their garage. The final velocity range we can tolerate for monitoring motion is 0.152 to 3.03 m/s. Note that some cheap LIDAR/ToF sensors can have repetition rates up to about 8 times faster.

The calculation above suggest that the HC-SR04 ultrasonic sensor, because of cost, the number of pings we can do within the time constraints, and at first glance, the range is a reasonable starting choice for measuring distance and motion. Similar range limit calculations need to be considered for any object, not just a car, and not just an ultrasonic sensor.

Motion monitoring

The discussion of velocity limits only gets part of the way in analyzing motion. The first rule of any motion detection scheme is: *All motion detection methods are based on difference measurements.* We found that an error of 0.5 cm, translates to a lower velocity limit of 15.2 cm/s. However, if the difference between two successive ping distance measurements is less than 1 cm, we cannot reliably indicate whether motion has occurred. Because each of the two measurements has the same error, if a first echo returns a distance of 10 cm difference, we can say the sensor to object distance will be between 9.5 cm and 10.5 cm. if the second ping returns an echo of 9.5 cm. should we really assume there was motion of -0.5 cm, when the error on that second measurement means the true value can be

somewhere between 9.0 and 10.0 cm, overlapping the range of the first measurement? However, if there is a 1 cm difference, we are beyond the error limits, and have a higher probability of assuming motion has occurred. Clearly, we must carefully analyze the error of our measurements. We will need to somehow adjust our data acquisition strategy. We have three options:

1. Change the ping interval time, so that two successive pings are above the error limits of the sensor, and encompass the anticipated object velocity;
2. If the minimum ping time interval is very fast, relative to the objects velocity, average many ping echoes at the maximum expected ping interval time, and compare the average values between intervals;
3. Monitor over a range of successive maximum ping intervals, so that the cumulative result of the difference measurements is above the velocity range being investigated.

The most flexible case, with the best motion detection possibilities, but also with the most complications, is option 3. It is that case that we will investigate in detail. Case 3 is about using an array to hold values. For efficiency, we want a special array - a circular buffer - and even more specifically, a first-in-first-out (FIFO) circular buffer. We then continuously analyze the data in the buffer.

Monitoring changes over an array can be as simple as a moving average, or as sophisticated as using a Kalman Filter. Average measurements of a buffer is usually the simplest method for monitoring the velocity that immediately comes to mind. We continuously compare new measurements to previously stored buffer averages. Averaging subsets of a stream of data is the basis of a moving average calculation. We define the moving average interval by adjusting the buffer size. We monitor changes (differences) in the moving average at longer time intervals, and compare to the rate of change we anticipate the object's motion will produce. If you are familiar with Control Charts for monitoring industrial processes, or market volatility, you are already familiar with the basic process. Moving averages are absolute time or distance values, not relative values. This makes the data dependent on position from the sensor.

Another method calculates the slope of either the distance or echo times in the buffer, derived from a least squares regression. Computationally, this might require more cpu time than moving average. An advantage in this case, is there is no need to store previous averages. Instead, we look for changes above some standard minimum slope that represent sensor noise slope limits. In addition, the method produces relative reference differences.

Both classes of measurements are lagging indicators of change. The more successive values we incorporate into our analysis, the further in the past we are looking at mobility changes. The closest we can come to "now" is the difference between two successive values at very small time differences. but we have to be aware of the limitations. This is part of the detailed discussion below.

I will concentrate on the slope method. However, a moving average method would have a similar overall process to achieve motion detection, and have similar statistical problems.

The size of buffer to use depends on the motion and static error of the measurements. The larger the buffer, the better the statistics related to deciphering the object's motion behavior, but increasing buffer size introduces lag in determining motion. This can be good or bad. As we increase the buffer size, we have more confidence in the difference in successive values, and minor velocity variations and sensor errors are minimized. At low object velocities, a larger buffer may be required to overcome statistical errors, but this extra time introduces a lag, limiting the idea of "real time" measurements.

The previous discussion sets the broad stage. The devil is in the details. A lot of what follows was developed over a lengthy trial and error period from an overly naive starting point, to many, "Ah ha", and "Oh shucks" moments. Most of the discussion is generic, but some is specific to the project.

As suggested, the starting point in any motion measurements is to establish the baseline for the errors in the sensor measurement. As previously noted, HC-SR04 specifications indicate the distance values are within 0.5 cm. What does that mean? Is that the average error, the variance, standard deviation (σ) or how many standard deviations, 2σ , 3σ ...? For static measurements, where 0.5 cm error may be negligible, or we have the luxury of improving the error by averaging multiple echoes, the error may not be a big deal. In slope measurements, we need a more thorough understanding of the statistical spread of sensor errors.

In principle, determining the error statistics is simple. Pick a target, place it a distance comparable to the maximum distance range you will be measuring, calculate the ping time interval based on the maximum distance to the target, add a few ms as a safety measure to avoid echo conflicts, and record a lot of measurements, say 2000 or more. Well, that works, but with all sorts of caveats. We will first look at temperature and target characteristics that impact every project using sound, then look at target/object issues.

Temperature and humidity and sound velocity

Sound velocity, or echo times are dependent on the temperature and humidity ranges under which you are making the measurements. Especially if you expect wide variations and desire fairly accurate measurements this may be critical. There are a number of online calculators to get the sound velocity as a function of temperature, humidity and sometimes barometric pressure. However, I found some discrepancies between calculators that appears to be related to the original data used to develop the calculation method. It appears one of the best references is:

The variation of the specific heat ratio and the speed of sound in air with temperature, pressure, humidity, and CO2 concentration. Cramer, Owen, *The Journal of the Acoustical Society of America* 93, 2510 (1993)

There are very good calculators based on this work. The original data fitting equation has 16 terms. I wanted to try and reduce that a bit, because for an ultrasonic sensor that kind of calculation is overkill, and I felt that someday I might need an equation that could be used with fewer cpu cycles. The next couple of paragraphs are a diversion from the main topic, but may be useful to some readers.

See the Excel file *Sound_velocity_T_RH_Calculator.xlsm* for details. A table of data was generated from Cramer's equation. The original calculation was a JavaScript routine scrapped from the web site:

<http://www.sengpielaudio.com/calculator-airpressure.htm>, "which used all 16 terms of the original equation with saturation vapour pressure taken from Davis, *Metrologia*, 29, p67, 1992,...and a mole fraction of carbon dioxide of 0.0004.". The JavaScript was translated into an Excel macro. With the air pressure held constant at 101.325, which is at STP, a table of sound velocities as a function of humidity and temperature was generated. Using Excel's optimization function, Solver, along with some educated guesses on appropriate equations, the following reduced equation was developed:

$$\text{speed (m/s)} = 330.8966 + 0.540567 * T = 0.001776 * T^2 + 0.009432 * Rh - 2.98489E-05 * Rh^2$$

where temperature, T, is in centigrade, and Rh, is the relative humidity, in %.

Errors in the range of 0-30 °C and 0-30 Rh are almost negligible for the kinds of measurements discussed here: In the most important ambient ranges, i.e., 15 - 30 °C, and 30 - 75% humidity, the errors are within +/-0.05 m/s, which is ~ 0.015 % (The specific values chosen for comparison have to do with the experimental conditions under which I did my own static ping tests.) In all of the cases, where I needed to calculate cm instead of μs , I used 0.034445 cm/ μs as the velocity of sound.

Attempts to downgrade the equation to use just the non quadratic terms, produced a slightly higher error, but if speed of calculation is critical, the equation below would still be acceptable in the region of 0-40 °C T and 0-40 Rh.

$$\text{speed(m/s)} = 330.3633 + 0.614639 * T + 0.006749 * Rh.$$

This formula turns out to be very similar to an equation often used to calculate sound velocity for the ultrasonic sensor, which does not explicitly consider the relative humidity. Whenever you do sound measurements, the temperature and relative humidity should be recorded. All your distance measurements depend on it. Getting the temperature and humidity may not always be possible, so you do the best you can. I happen to have a certified temperature and humidity meter. At the very least, measure the temperature, and assume a humidity based on the latest weather report, unless you are in an air condition environment, where the humidity may likely be under 50%. Even then, depending on the room conditions and how long the measurements take, both the temperature and humidity will likely rise a bit. I monitored the values at the beginning and end of a run, and averaged them if necessary. In the current case, sound velocity was 0.034445 cm/μs, or 34.44 cm/ms, at 22.2 °C and 43 % Rh.

Target Characteristics

A target that is very similar in material, shape, and size of the intended target should be used, because as we will see, the material will have an impact on the observed standard deviations. Different materials have different sound reflectance. Metal, human skin, water, plastics, foam, cloths will be very different. They can absorb negligible sound energy, such as highly polished metal surfaces, or absorb almost all the sound energy, like open weave cloth. Foam can have a range of reflections: from absorbing most of the sound energy to reflecting most of it, depending on type of foam cell and surface coating.

The shape of the surface is important. If the material is not flat, or is not normal to the ultrasound sensor centerline, depending on the curvature and angle to the sensor, some or none of the reflected sound will reach the microphone.

Surface texture is important. A rough surface (diffuse reflection) versus highly polished surface (specular reflection) of the same material will impact the sound energy returned to the sensor. That is why it is so important to try and use a target with similar properties.

Admittedly, the cases I investigated below violate most of these rules at some level. My first measurements were not from the front surface of a car; I had to start with something simple to understand the echo dynamics.

Static measurements

As already clear, the basis for all motion calculations stems from error statistics derived from static measurements. For static measurements, I started with a 26"W x 69"H x 1" foam insulation sheet (Owens Corning, F06) with a very smooth, flat surface. The sheet was carefully mounted perpendicular to the HC-SR04 sensor, vertically and horizontally, at a distance of 2.66 meters, limited by the size of the room and furniture. The sensor was aimed at the center of the sheet. Two small sketches were used to obtain the static data:

```
/*
  usonic_static_2000.ino
  Uses the primary examples from the author of NewPing to drive sonic sensor:
  HC-SR04 NewPing Iteration Demonstration
  HC-SR04-NewPing-Iteration.ino
  Displays results on Serial Monitor
  */

// Arduino/sensor set up
// Uses Serial Monitor to display Range Finder distance readings
// Include NewPing Library for HC-SR04 Ultrasonic Range Finder
#include <NewPing.h>

// Hook up HC-SR04 with Trig to Arduino Pin 4, Echo to Arduino pin 5
// Maximum Distance is 400 cm (4m)

#define TRIGGER_PIN 4
```

```

#define ECHO_PIN 5
#define MAX_DISTANCE 400

//Initialize the sensor pin setup and distance
NewPing sonar(TRIGGER_PIN, ECHO_PIN, MAX_DISTANCE);
unsigned int duration;
float distance;
const int iterations = 2000; //must change this to 500 if Uno, because cannot handle greater array size.
//unsigned int pings[iterations]; uncomment this if dumping to array before printing.
int cnt = 0;
void setup() {
  Serial.begin(9600);
  //If dumping to array use the following loop; else comment out this entire loop out, and
  //uncomment next loop
  for (cnt = 0; cnt <= iterations; cnt++){
    duration = sonar.ping(); //single pings @delay time.
    distance = (duration / 2) * 0.03445;
    //pings[cnt] = distance; //uncomment if using array to hold distance values before printing.
    // if using array to hold data, comment out the next 3 lines
    Serial.print(duration);
    Serial.print(", ");
    Serial.println(distance);
    delay(33);
  }
  /*
  //uncomment this block to test whether Serial causing issues while blocking the Serial.print lines in the previous loop.
  //TURNS OUT TO HAVE NO EFFECT ON OUTPUT VALUES.
  for (cnt = 0; cnt <= iterations; cnt++){
    Serial.print(pings[cnt]);
    distance = (pings[cnt] / 2) * 0.034445;
    Serial.print(", ");
    Serial.println(distance);
    delay(33);
  }
  */
}
// loop() is not used, because only want one pass.
void loop() {
}

```

This sketch, could be run with either immediate printing, or with dumping to an array, then printing. Tests showed the results were the same in both cases, but the latter was limited by the number of pings/run, and had to be run several times to match the same ping count as for immediate printing.

A second sketch which produced the same range of values within experimental error, was also used to ensure the data were not affected by printing to the Arduino IDE Serial Monitor, but had to be run at much lower pings per run because of array size storage issues. That sketch directly accessed the Uno time registers. The data was then printed to the Serial Monitor in a separate loop. The reference for this is:

<https://www.davidpilling.com/wiki/index.php/HCSR04#what>

The only reason for running different sketch variations was my unfamiliarity with my Arduino Uno and ultrasonic sensor. I had to assure myself of just how fast the Arduino could process the information.

Nearly all the data were analyzed using MS Excel 2007. I can reasonably quickly write Excel vba macros, to manipulate and plot data, which makes Excel particularly useful for me. If you are familiar with Open Office, that will work too, unfortunately with a different macro language. Many of the more computationally involved calculations or just drudge calculations were done using macro methods. The Excel file is in the current repository: *minimum_ping_cycles_for_distances.xlsm*.

For some readers, who need more detail, this Excel file may not be as straightforward to navigate as you would like. It contains all the data and most of the codes, except for Autocorrelation and LOESS curve fitting. Some effort and head scratching will be necessary to fully understand some spreadsheets. The file was always a work-in-progress, and not written for casual public consumption. There are methods, backup tests and some failed ideas represented in some of the sheets. The macro codes are not optimized or pretty, but they work. For anyone wishing to really understand what I did, having the original data may offer deeper insight, or deeper criticisms. Definitely read the text blocks associated with each worksheet, they should help decipher the intentions of the spreadsheet. Several worksheets, especially those that calculate buffer parameters may be quite useful. In this discussion, unless a specific file is specified, the format, *[worksheet_name]* represents a single spreadsheet in the above referenced Excel file.

All the static, streaming echo data was copied to the Excel spreadsheet from the Serial monitor by simply highlighting all the values on the Arduino IDE Serial Monitor, pressing Ctrl-C, copying to a worksheet, and using the Excel *Data/Text to Column...* process to split the data text lines to separate cells. The average value and the standard deviation were then calculated, and all the data plotted as a function of a running integer index. The index can be considered to be a normalized time index, because all data used a 33 ms ping interval.

To better understand the static error distribution, Excel's Histogram function was used (Microsoft Data Analysis Pack). The naive belief was that the data would be normally distributed. It was not.

The plot below shows the frequency counts and cumulative data.




Figure 1. Frequency plot of microsecond echo times at 1 μ s intervals, for Foam Insulation.

Initially, this data was confusing. Only after some in-depth online reading, did I come to believe I understood this data. Three articles turned out to be important to my comprehension.

Particularly, the first article below was very helpful, but did reach beyond my understanding of electronic circuitry:

[Making a better HC-SR04 Echo Locator](https://uglyduck.vajn.icu/ep/archive/2014/01/Making_a_better_HC_SR04_Echo_Locator.html)

https://uglyduck.vajn.icu/ep/archive/2014/01/Making_a_better_HC_SR04_Echo_Locator.html

[Why ultrasonic module sends out 8 cycles ? and , why the trigger pulse is 10us?](https://electronics.stackexchange.com/questions/274249/why-ultrasonic-module-sends-out-8-cycles-and-why-the-trigger-pulse-is-10us)

<https://electronics.stackexchange.com/questions/274249/why-ultrasonic-module-sends-out-8-cycles-and-why-the-trigger-pulse-is-10us>

[HC-SR04](https://www.davidpilling.com/wiki/index.php/HCSR04)

<https://www.davidpilling.com/wiki/index.php/HCSR04>

The frequency plot shows that even though the data was binned at 1 μs intervals, there are many empty bins. The frequency plot distribution, and especially the cumulative data plot, identify eight "distribution envelopes", each envelope with 6 clear regions within.

What is going on here? Based on the cited information and my own ideas: There is a 10 μs trigger pulse, followed by a 248 μs delay, while the sensor powers up or down, followed by the 8 pulse emit cycle, lasting just under a total of 200 μs , then the sensor waits for echoes. The series of 8 pulses as a 40 KHz signal out, means that each pulse is 25 μs long.

In a perfect and naive world, the first pulse sent from the sensor is the first pulse back, which triggers the sensor to put the echo pin high, triggering the Arduino to query it's timer. The ATmega328 timer that the NewPing library uses has 4 μs resolution, so when the sensor triggers the echo pin high, the Arduino responds with a timer value on the next 4 μs pulse, and gets the time difference from when the trigger pulse was sent to the sensor. That explains why the data counts are multiples of 4 bins.

As for the eight envelopes, it seems clear they represent the eight pulses. The problem is that triggering is not perfectly linear, it depends on electronic resonance factors in the circuitry, so that depending on the sound energy returned to the sensor, triggering may occur on any of the roughly 6, 4 μs timer values that intersect the returned pulse and trigger the sensor to drive the echo pin high, (about 25 μs envelope width). The returned sound energy may not have enough energy to trigger the first pulse, and gets triggered by the second or third 25 μs pulses, etc, again with a resolution of 4 μs . What we end up with, is what you see in the frequency binned data.

There is some interesting stuff going on in the sensor circuitry from this data. At the particular distance, and with the insulation foam target, the first pulse triggers less than 4% of the time. The second pulse triggers the echo pin about 38% of the time, and the third pulse about 32% of the time.

As a comparison, a small aluminum panel was also run under static condition, albeit at a closer distance. The plot below shows a similar, but more truncated pattern.




Figure 2. Frequency plot of microsecond echoes, at 1 μs intervals, for aluminum sheet.

Instead of a clear indication of eight pulses, only 4 pulses are evident, each envelope still with 6, 4 μs sub counts. The reason for the tighter distribution is likely the combination of the closer distance to the sensor and the more reflective metal. The foam insulation target, even though it had a hard foam surface, likely absorbed some of the sound energy compared to the aluminum panel and was further away. The different sizes of the material had no

bearing on the data. Prior beam spread sensor tests showed that at the distances measured, the entire ultrasonic beam would have intersected the panels. In the metal target case, the first ping triggers the echo pin a little under 10% of the time.

Although the data is not presented here (it is in the Excel file), experiments were performed with the foam target at closer distances to the sensor, but with only 500 pings per run. At the shortest distance, the data do show fewer envelopes, similar to the metal target case. Thus, it is important, especially when dealing with motion measurements, to understand the target's material, shape, and distance from sensor conditions, and of course any major changes in temperature and humidity.

The data highlight several issues that impact how we ultimately use the information to determine motion. The most common standard deviation formula we rely on assumes data is random and normally distributed. Clearly these distribution are highly skewed, no matter what the material or distance. If we close one eye and squint with the other to assume a normal distribution, the standard deviation is $32.8\ \mu\text{s}$, which translates to $0.56\ \text{cm}$. So the specification error typically listed, of $0.5\ \text{cm}$, for at least the sensor used here, would appear to be close to the standard deviation, but the skewness of the data means we should be wary in our error estimates and other calculations.

Later, we will be more interested in the 95% tolerance or prediction interval, which is $64.2\ \mu\text{s}$ ($1.10\ \text{cm}$). Despite the concern that this is not a normal distribution, we can get the 95% value by manual inspection of the cumulative data, and looking where each "wing" beyond the main peak shows 2.5%. What we find is the difference is $126\ \mu\text{s}$, or $63\ \mu\text{s}$ from the median. This is close enough for our purposes to suggest we can get away with just assuming the common standard deviation is a "good enough" estimate. Note that the mode is $15447\ \mu\text{s}$ and the mean is $15472.8\ \mu\text{s}$ in the foam target data, which represents quite a large shift.

Another statistical roadblock is that the raw histogram data clearly show that certain echo bins are "favored" or more prominent than others. As already suggested, this is not a sampling issue, but reflects certain biases in the sensor and Arduino circuitry. This means no distribution, no matter how many pings we do, will ever follow a smooth monotonically increasing and decreasing envelope. Certain distribution bins will always exhibit "favored" values. In addition, tests done at different distances (in Excel file, *[foam_multiple_runs]*) indicate that the standard deviation is distance dependent, because of the triggering dynamics.

In all subsequent discussion, the foam insulation data is used. It makes more sense to work with a worst case to develop a robust process, than the best case, unless experiments clearly provide a reason to do otherwise.

Nearly exclusively, I will discuss monitoring motion in terms of sensor echo time per time interval, such as $\mu\text{s/s}$ or $\mu\text{s/ms}$, not distance/s or distance/ms. With NewPing, the raw data from the Arduino is in microsecond (μs) differences. The time difference to distance calculation just wastes cpu cycles. Using just the time values removes one degree of potential interpretation issues.

Car Velocity and Statistics

The anticipated maximum distance, the objects velocity, and the statistical error limits control the boundaries of motion measurements. Here is a real example to highlight the considerations. First, we need some idea of the motion of the car in the range we are going to monitor velocity. We want to trigger a result or event when a car pulls into a garage. We want to be sure that at least 95% of the time or 1 in 20 measurements, we can reliably say the car is moving. As we discussed earlier, the total time frame of the measurements is around 11 seconds to negotiate $5.5\ \text{m}$. We are going to consider constant velocity, although it is definitely not true throughout the distance range. After a bit of units manipulation, the car's velocity is $0.049\ \text{cm/ms}$.

From static experiments on our sensor, we know that the 95% of the static values fall within an echo value range of $\pm 64 \mu s$ (prediction limit). Beyond that limit, there is a only a 1 in 20 chance we will be wrong in our prediction that the car was moving (really a bit less). A difference of $64 \mu s$, using a sound velocity of $0.034445 \text{ cm}/\mu s$ corresponds to a distance of 2.2 cm. However, that is the time for sound to get to the car and back to the sensor, so the true distance difference is 1.1 cm. We decided at the outset that we would start monitoring the car at 18 feet, which is 5.5 meters. The time needed for the sensor to send a signal and return an echo is, $2 \times (5,500 \text{ cm}) / (0.034445 \text{ cm}/\mu s) = 31,856 \mu s$, or 31.8 ms. This defines the very minimum time difference between pings. The minimum car velocity that will meet the criteria to reliably say it will move 1.1 cm in 31.8 ms is $1.1 \text{ cm} / 31.8 \text{ ms} = 0.0348 \text{ cm/ms}$. We compare that to the average car velocity we calculated of 0.049 cm/ms ; this value is higher than the minimum velocity we just calculated. so a 33 ms ping works to detect the moving car, with a 5% chance we will be wrong (actually a bit better, because 5% is the remaining counts on both sides of the distribution).

We can put this all together in a single formula for quick calculation:

$$v_{min} = \frac{890 \times \sigma_{static} \times v_s}{t_I}$$

v_{min} in cm/ms, is the minimum velocity of an object to reliably (95% certainty) indicate motion;

σ_{static} in μs , is the standard deviation of the static measurements;

v_s in $\text{cm}/\mu s$, is the velocity of sound, either calculated or estimated as $0.034445 \text{ cm}/\mu s$;

t_I in milliseconds (ms), is the interval time between pings;

890 is the product of 1.96 (2σ factor) \times 0.5 (ping to object and back) \times 1000 (convert μs to ms)

This is compared to the expected velocity of the object (in cm/ms). [WATCH UNITS!] If $v_{expected}/v_{min}$ is greater than one, the ping interval is adequate to indicate motion, if less than one, either increase the ping interval time, or add more ping intervals. More, later on this point.

Not all cases will have this kind of negotiation between ping interval and object velocity. Obviously, if object motion is very slow, we have the luxury of averaging two bunches of ping measurements at separate large ping intervals. Subtracting the results, we can get a fairly good idea if motion has occurred. When should we worry about the tradeoff between ping interval and object velocity? It has to do with confidence intervals, which are related to the standard deviation, and provide us with some idea of whether we can assume a difference measurement has any merit.

To demonstrate why this averaging is useful, we are going to use real ping values derived from the static ping measurements shown for the foam target that were discussed previously. For a 2 buffer array, we found the standard deviation (σ) to be $21.2 \mu s$. This is the standard deviation from analyzing 1,999, 2-element buffers. We found σ to be $9.4 \mu s$, for 1,991, 10-element buffers. We want to be certain of our measurements with a 95% confidence limit. Confidence intervals are not the same as the prediction level. The former tell us how close to the true mean an experimental mean will likely be; The more points we gather, the more likely the experimental mean will be closer to the true mean. This is also the accuracy. Confidence intervals therefore depend on the sampling size. Prediction level tells us the expectation that a value is within the distribution, i.e., whether the slope is due to static or motion changes; this is related to the precision of the measurements. As the confidence interval increases, the precision goes down.

Statistics tells us we calculate the 95% confidence intervals for these two cases as $1.96 \times \sigma / \sqrt{N}$ (assuming normal distributions). The values for the 5 element and 10 element buffers are $29.4 \mu s$, and $5.8 \mu s$, respectively. If we assume the average of either the 2 element or 10 element buffer was $15400 \mu s$, the confidence interval for the 2 element buffer tells us that we can expect the true value to be somewhere between 15330.6 and 15429.4 , and for the 10

element buffer to be between 15394.2 and 15405.8 μ s. That is a big range difference between the 2 buffers. If we could afford the time to average and subtract two, 10 element buffer arrays. we could attribute any difference over $\sim 11.6 \mu$ s to motion, which corresponds to a distance variation of 0.19 cm. If we only considered a two element buffer, the difference turns out to be 58.4 μ s, i.e., 1 cm, 5 time higher. (The confidence limit was doubled, because two readings have the same confidence interval, so to make sure they potentially overlap by 5%, we have to double the confidence interval.)

A definite complication is that the car velocity is not likely to be a constant 49 cm/s as in our example, unless we typically race into the garage and stomp on the brakes, just before we hit the wall. According to the calculation, we do have a little bit of extra margin, but very likely it will not be enough. So that suggests we cannot trust simple time differences at 32.8 ms ping intervals to monitor the distance. If any measurement drops below that limit, the Arduino would stop responding. The goal is to have an LED strip that changes color when the car is moving in the garage to indicate, both how close we are to a designated stop position. and the lateral car position. It will not do, if the LEDs go black because we have slowed below the expected velocity. A buffer, because it is a lagging indicator, can alleviate this problem to an extent, by smoothing out rapid changes.

Using the Slope of Buffer elements to Monitor Motion

As indicated in the Introduction, calculating the slope of echo readings in a buffer was the choice for monitoring motion changes. The preceding section, dealt with raw static streaming measurements and statistics, which are the basis from which we extend to the true level of monitoring motion, by calculating the slope of the buffer after every ping.

As with any lagging indicator, the advantage and disadvantage of monitoring the slope of a buffer is a balancing act. The advantage is a smoothing of the data from rapid single measurement fluctuations because of higher confidence limits; the disadvantage is the data lags behind what is really happening right "now". Depending on how we set the minimum limit to trigger a motion event, and the ping intervals, a larger buffer will "slide" larger echo time changes to the front of the buffer, as the buffer fills, so even if the car velocity begins to drop, we will still see an impact on the slope of the earlier changes.

The size of the buffer influences the limiting slope beyond which we can say motion has occurred. We need to find the limiting slope for a certain buffer size. We can take the stream of static measurements we made (~ 2000) and get the slopes by moving a buffer "frame" of some size along every echo of the stream of echo data. The slope is calculated as the linear regression of the frame of N microsecond echo returns (y values), against a set of corresponding time "x" values 0, 1,...(N-1) as integer time. The "x" values represent a normalized time index, where the true time intervals, in this case, 0, 33, 66 ms...etc., are divided by the ping interval time, 33 ms, to provide a simple common integer index. This does not affect the calculation of the slope, as long as we are always considering the same time interval. If we started with 2000 static echo measurements, we end up with 1991 calculated slopes. We then calculate the σ and 2σ (95% prediction level) of the slopes. Below is the histogram/frequency plot of the linear regression slopes for the case of 1991, 10 element buffers, based on the static measurements:

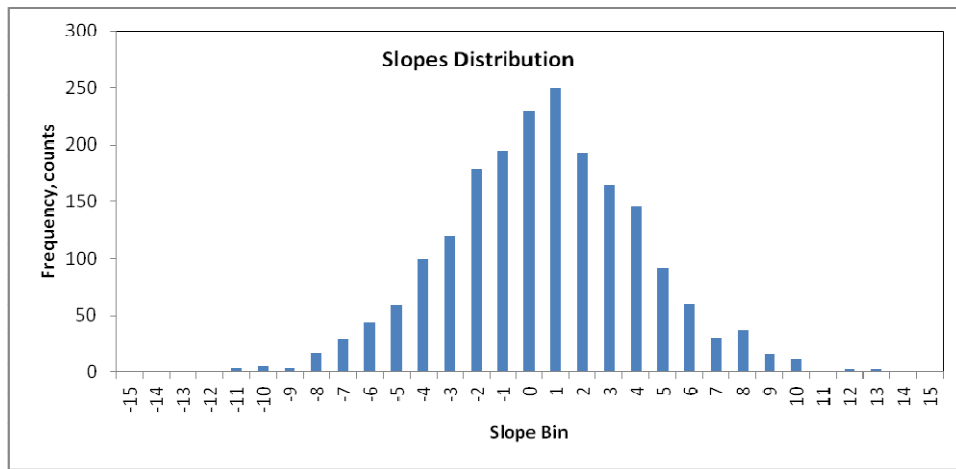


Figure 3. Frequency distributions of slopes for running 10 element slope buffer, at unit slope bins.

Why is there a distribution of slopes? The slope of a static measurement should just be zero? Only in a perfect world, with zero error. If you calculate the average slope of 2000 echo measurements, the slope is nearly zero (mine was 0.0068). However, a buffer is a far lower subset of echoes with some random error value between each echo. As suggested by the previous discussion of the 5 and 10 element buffers of just the static echoes, the confidence interval is a lot wider for these much smaller number of echoes. By analyzing the distribution of the buffer size slopes over all the static echo values, we can determine a slope limit, above which we can assume that we have motion.

We determine the slope limit from the standard deviation, to get the σ , and 2σ for all the slopes. We will only use the 2σ data, which tells us that only 5% of the slopes (1 out of 20) will fall outside the 2σ value. For this specific static case and sensor, the 2σ case was -7.08. A Q-Q plot still shows the data is not normally distributed in concert with the original data from which it is derived. This suggests that if we wish to detect motion using a 10 element buffer, we need to see a slope greater than 7.1 or less than -7.1, because the distribution is nearly symmetrical.

Some readers may be wondering why use a linear regression relationship as a buffer monitor when, with the exception of all buffer elements indicating a constant velocity change, or no change, the data should never be linear. There is no fundamental reason why other linear regression equations to determine a slope cutoff could be used. It is just the most convenient, because it provides single parameter to judge motion.

Is there an Optimum Buffer Size?

So as we now see, there are several variables in play in reliably determining motion of a object using a buffer: the static/slope error distribution, the size of the buffer, which affects the confidence interval of the buffer, and the ping interval. I showed above how to calculate, for any buffer size, the slope change limit we have to reach to indicate motion.

So is there some optimum buffer size to use, based on the velocity and ping interval? The short answer is it depends on what your goal is. The question comes down to asking how much time can you afford to recognize a change in object position?

Let us look at a hypothetical example of what happens to the slope distribution as a 10 element buffer registers continuous motion through the last three buffer elements. Remember, this is a FIFO buffer array, so it fills from the last element (9) to the first element (0). There are three cases: The calculation process is the following: For each element in the static array of 2000 recorded echoes, we take the first 10 echo elements, subtract some time value, in this case, we will use $-45 \mu\text{s}$, from the 9th static echo value element; this represents a movement toward the sensor with a target that will move -0.78 cm in the time interval, about 38% above the 0.56 cm 2σ value. We calculate and store the slope of the entire 10 element buffer chunk. For a decrease in microseconds over two elements, we take

the static value in the 8th element, subtract $-45\ \mu\text{s}$, and from the 9th element static value we subtract $-90\ \mu\text{s}$, which indicates that the object moved at constant velocity toward the sensor for two interval; again calculate and store the slope of the buffer. For the third case, we do similarly for 7th through 9th buffer elements; We then move down one static ping element in the stream and go through the same process again, until we have gone through 1991 static echoes of the 2000 we started with. We now have three slope distributions, each with 1991 points. We again do a frequency plot to show how the slopes change. The plot below shows the data.

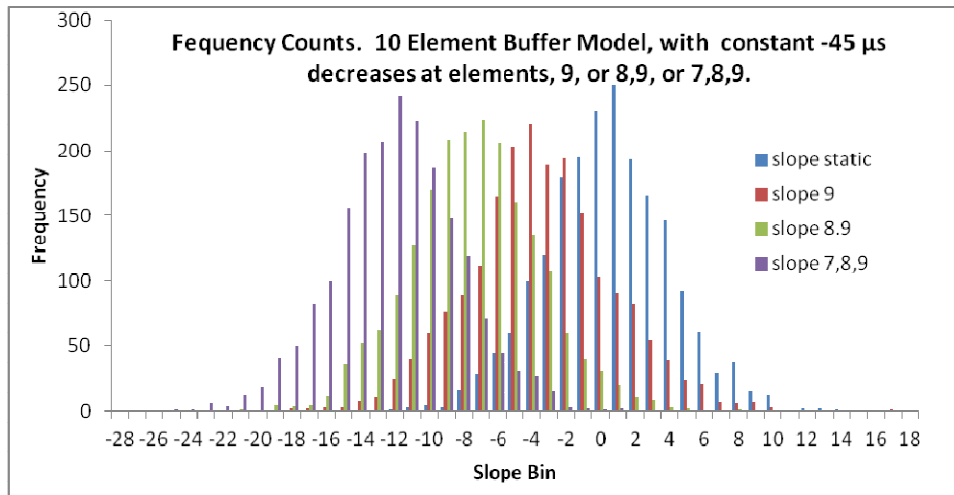


Figure 4. Frequency Distribution of static echo values modified with $-45\ \mu\text{s}$ drop (vel., $1.6\ \text{cm/s}$) at $33\ \text{ms}$ ping intervals

This plot is busy, but gives a visual idea of what happens to the slope distribution as an object moves toward the sensor. The blue distribution is the distribution of static slopes vales without any motion, the slope 9 (red) distribution represents a $-45\ \mu\text{s}$ drop in the 9th element. Note the large amount of overlap with the static slope distribution, which suggest for this case trying to determine a moving object with just the difference between two pings is not statistically a good idea. The green (slope 8,9) distribution represents roughly the case were we have a $-45\ \mu\text{s}$ drop between the 7th and 8th element, and an additional, $-45\ \mu\text{s}$ drop between the 8th and 9th element. This is starting to look statistically better. Finally, the last case where three buffer elements register a constant motion, shows that we have a good chance of deciding if an object is moving, for this specific case. Although there is still overlap, it is reduced to about 5% overlap with the static (blue) distribution. Note where the slope 7,8,9 distribution overlaps with the static distribution. It is around a slope of -7 .

We already showed how to obtain the minimum slope that different sized buffers must meet to indicate motion at least at the 95% confidence limit. What might be useful is to translate that minimum slope information, for each buffer size, to microseconds between buffer intervals, to indicate the velocity range a buffer can handle, and when the buffer would indicate motion has occurred. The calculation can be done manually, using trial and error or "Monte Carlo" methods, but is overly time consuming. Using the magic of computer optimization methods is a better approach.

I will mention two optimization methods: Excel's Solver optimization function, and a python module based on the SciPy library. There are many optimizers around; look into the statistical package "R", for many more, as well as SciPy. However, Excel's Solver optimization was the first method used to establish the correlations between buffer size, ping interval, and object velocity. Two worksheets [*theo slope & stdev*] and [*buff_limits*] and associated macros in the Excel file, make quick work of the problem.

For those without Excel, a python module, *bufferlimits.py* is included in the repository. The module will run on python 3+ versions. As input, it requires a csv file of a stream of static echo values, one echo per row. The output will be the same as the tables discussed below except as a csv file. Check the settings section of the module to get the desired buffer matrix. There are ways to directly log Arduino data to a csv file, which I later opted to do, but did not use for this calculation (see below).

The calculations proceed in the following manner: I described previously how the 2σ values of the slopes of each buffer size represent the minimum slope values, and how they were obtained for different buffer sizes by consecutively finding the slopes of chunks of data, equivalent to the number of buffer elements from the stream of static ping echoes. Worksheet [*theo slope & stdev*]. To get the minimum ping values per element in a buffer, we start by assuming a constant velocity for the object, and that the object was initially at rest, which means all the buffer elements are initialized to zero. We create N - 1 buffers of the size N we want. In the first buffer, we add some rough value x to the last buffer position, N, which we estimate from experience may be close to what may be required to cause the slope of the buffer to match the anticipated 2σ value. Continuing, in the second buffer, we add a rough estimate of the anticipated microsecond echo difference to the N-1 element, call it x', and add 2*x' to the Nth element. We continue this bumping process in the subsequent buffers. The trick is to minimize each of these N-1 buffers separately, so they match the 2σ minimum static slope we found (-7.08) . This is where Solver comes in. The result is a matrix with rows representing the buffer elements, and the columns representing the minimum echo value "jumps" between buffer elements to meet the requirement that motion has been detected. For a calculated buffer slope to indicate a moving object, every ping echo must meet or exceed, the ping times shown in the buffer. (The equation for starting numbers, representing "experience" can be found in the coding; euphemistically, it's "derivation" is shown in the Excel file [*buff_limits*].)

Below is an example of the data produced for a 10 element buffer.

Ten element buffer with optimized μ s values, Minimum slope = -7.08; Ping interval = 33 ms.

Buffer Number									
Buffer elem.	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	-7.08
2	0	0	0	0	0	0	0	-7.49	-14.15
3	0	0	0	0	0	0	-8.34	-14.97	-21.23
4	0	0	0	0	0	-9.81	-16.68	-22.46	-28.31
5	0	0	0	0	-12.29	-19.63	-25.02	-29.94	-35.39
6	0	0	0	-16.68	-24.58	-29.44	-33.36	-37.43	-42.46
7	0	0	-25.39	-33.36	-36.88	-39.25	-41.71	-44.91	-49.54
8	0	-46.71	-50.77	-50.05	-49.17	-49.07	-50.05	-52.40	-56.62
9	-129.75	-93.42	-76.16	-66.73	-61.46	-58.88	-58.39	-59.88	-63.70
vel us/ms	-3.93	-1.42	-0.77	-0.51	-0.37	-0.30	-0.25	-0.23	-0.21
vel cm/s	-67.72	-24.38	-13.25	-8.71	-6.42	-5.12	-4.35	-3.91	-3.69
lag factor	33	66	99	132	165	198	231	264	297

There is a lot of information in this table:

1. The first non zero value in each column provides the limiting microsecond echo change or "jump" between elements needed to show 19 out of 20 times that motion has occurred by that element. Of course, this difference must be at least as great for the later buffer elements, as well.
2. The last value in a column represents the total change we need to see in the buffer to declare motion.
3. The magnitude of the jumps from buffer to buffer (first non zero value in each column) is a measure of how the confidence limits change based on the standard deviation, and buffer size. Confidence intervals are not the same as the prediction level. The former tell us how close to the true value an average will be; prediction level tells us the expectation whether the value is within a 95% envelope or not, due to static or motion changes.
4. The table also shows the limiting velocity that an object must have to trigger motion detected, with the number of non zero elements filled. The velocity, in contrast to the microsecond jumps, depends on the ping interval, and is calculated from: $velocity = 0.5 * echo_difference / ping_interval * sound_velocity * 1000$ where *echo_difference* is the first non zero element in a buffer, and *ping_interval* is the delay time between ultrasonic sensor triggers.
5. Lastly, the delay time is shown that will be needed to insure motion at the given velocity and ping interval.

As an example, consider the second buffer column in the table. If you want to know if motion is occurring after three pings [0, -46.71, -93.42] the last two echoes must exceed a -47 μ s difference each to register object movement. Given a 33 ms ping interval, the object must be moving at a minimum of 24.4 cm/s, or 9.6 inches/s to register motion. The delay time to register this motion will be 66 ms, which means we have to do at least three pings. If we expect the object will move much slower, and we are limited to the same ping interval, we look to buffers to the right. The lowest velocity we can monitor with a 10 element buffer at 33 ms ping interval is 3.7 cm/s (1.4 in/s), but the delay time to detect that motion is nearly 300 ms, requiring 9 pings. You can now see why the artificial test case was chosen as -45 μ s jumps between buffer elements. Only in the third column do we see that with three pings, all above 25.4 μ s, can we say that motion occurred.

Let's compare to a five element buffer, with same ping interval.

**Five element buffer with optimized μ s values, Minimum slope = -19.78;
Ping interval = 33 ms.**

	Buffer			
buf. elem.	0	1	2	3
0	0	0	0	0
1	0	0	0	-19.79
2	0	0	-24.73	-39.58
3	0	-39.58	-49.47	-59.36
4	-98.94	-79.15	-74.20	-79.15
vel us/ms	-3.00	-1.20	-0.75	-0.60
vel cm/s	-51.64	-20.65	-12.91	-10.33
lag factor	33	66	99	132

What can we conclude between the these two cases?

1. Of course, the limiting slope of the 5 element buffer is higher because of the reduced confidence interval of 5 versus 10 echoes.

2. A smaller buffer does not have the range of velocities that a larger buffer can accommodate; this is a consequence of the lower statistical confidence levels of the smaller buffer. A larger buffer is more sensitive to lower velocities, and will therefore be less sensitive to changes in velocity.
3. Given the same number of pings, the smaller buffer more rapidly detects motion at a lower velocity than the larger buffer, i.e., the lag factor is lower.

Non Constant Velocity

The above discussion is based primarily on the assumption of a constant velocity. In many cases, the velocity will be constant over the short time intervals over which a buffer is filled. For example, a 5-element buffer fills in 165 ms, and a 10 element buffer in 330 ms. How about a case in which the velocity changes on a time scale within those ranges? There are an infinite number of variations of this case, but maybe some general comments might help.

If the velocity varies a lot around an average, some buffer element values can be lower, as well as other elements higher. The slope "averages" these changes, but only if enough points are being used to detect motion. Only specific trial and error can indicate what buffer size and ping interval will work using the slope detection method.

We can show one case that might provide some perspective on motion sensitivity. We will take the specific case of a 10 element buffer and a 5-element buffer, where motion occurs over two pings, i.e., the last two elements of the buffer show a decrease from the static data, and then the object stops. Of course the buffer continues to fill, and the elements that indicate the change are pushed "up" and ultimately out of the buffer. The plots below show the slopes of the traveling buffer as it encounters the motion spike. Remember, we are calculating the slope after each ping. The velocity of the object is set at $-46 \mu\text{s}$ per ping interval of 33ms. The two plots below represent the case of a 10-element buffer that starts filling at ping number 16, and a 5-element buffer that starts filling at ping 21. This is done so that the buffers both have encountered the two $-46 \mu\text{s}/33 \text{ ms}$ successive increases at the last two pings of the buffer.




Figure 5. 10 element buffer: Static object that suddenly starts to move for 66 ms then stop. Buffer slope changes from static slope values.




Figure 6. 5 element buffer: Static object suddenly start to move for 66 ms. Buffer slope changes from static slope values.

The way this data is generated is to randomly pick a region of the stream, and add to the original static echo returns the anticipated velocity to two elements in sequence, and then stop the addition (no motion) for further pings. In the cases shown, the streaming echo region chosen presents one of the worst cases. Note that the slope at ping number 19 is outside the 95% envelope. It is a single outlier. The vertical axis is the slope. The *slope limit line* (green dash) is the 95 % confidence interval; *static* (blue) is the slope calculated from 5 or 10 successive buffer points; *slope* (red) represents the case where the last two buffer elements at the start ping interval indicated, experience two successive $-46 \mu\text{s}$ drops, representing an object approaching the sensor at 30 cm/s. The motion then stops, but pinging continues.

Comparing the two plots you can clearly observe the much less noisier slope value in the 10 element buffer, note that the slope scale of the 5-element case is even compressed relative to the 10-element case. As expected all values, but the outlier are within the 95% envelope. The regions that standout in red are where the "motion" was instituted. While both cases show an effect of the outlier spike, the 5-element buffer, shows a much less definitive motion trigger. Three of five echoes are definitive for motion, and the last ping slope is at the borderline, while the 10 element case shows 7/10 motion trigger events. The outlier spike is evident, but it's effect has been "smeared" out better than in the 5 element case, because of the better statistics for a 10 element buffer. However, because we adjust the buffer slope limit based on the buffer size, either buffer would register that motion is occurring.

All the above shows what will happen at an assumed, and somewhat arbitrarily chosen buffer size and ping interval. These two parameters dictate what velocity range a specific buffer can handle. Can we work the other way around and get an idea of what buffer size might be appropriate to use based on the velocity of the object, and the ping interval? The Excel file *[calc_buffer_size]* or a python file, *buffer_size.py*, can provide the information. The buffer size that is calculated assumes that every element in the buffer is changing monotonically, representing constant object velocity throughout the entire buffer time length. In that sense, the buffer that is recommended is a lower limit, especially if there is a desire to predict motion in only a partially filled buffer. At that point, shifting to *[buff_limits]* or the python calculation, which develops the correlation of buffer size with object velocity can be more useful to zero in on a final buffer.

In all the discussion above, a ping interval of 33 ms was used. As discussed early on, this interval takes into account return echoes from objects or other reflective surfaces at about 5.7 m, or 18.6 ft. Most of the examples using NewPing, use a ping interval of 33 ms between pings, because it is near the sensor detection limit of the return echo. It is possible to use lower ping intervals, but only if it can be ascertained that pings will only occur from the target object. Sound spread from an ultrasonic sensor does not spread out in the typical radial pattern like most sound waves are represented. The reason is that the sound pressure is not transmitted uniformly. See the "What Went Wrong" section below for more detail on this issue.

Enhanced Detection

At this point, it is possible I have confused and lost a substantial portion of less dedicated readers, but there is an additional layer of analysis of buffer data that I will just touch on. The details are in the Excel file. Linear regression is used to determine the slope. A buffer is only well correlated to linear regression, when there is no motion detected in any buffer elements, or when all buffer elements indicate constant motion. All other cases produce lower slopes, with the data points either in a concave or convex pattern with respect to a linear slope line. Especially, in cases where we are using a large buffer, there might be cases where it is useful to know whether the motion is just being detected, or has ended, and is waning. If we calculate the quadratic fit to the buffer points ($y = ax^2 + bx + c$), where x is buffer element number, and y is the data echo value, the second derivative of the equation with respect to x , which is $2a$, tells us whether the data is concave up or concave down. If negative, the curve is concave downward (like a rainbow); if positive, it is concave upward. In our case, a negative 2nd derivative value means that motion is just starting to show in the buffer (elements 7-9). Once again, we have to define lower statistical limits to the value of the 2nd derivative to distinguish it from the buffer. I will just show a plot of a model case where, were the first four or last four elements of a 10 element buffer represent motion, and plot the 2nd derivative against the linear regression slope.

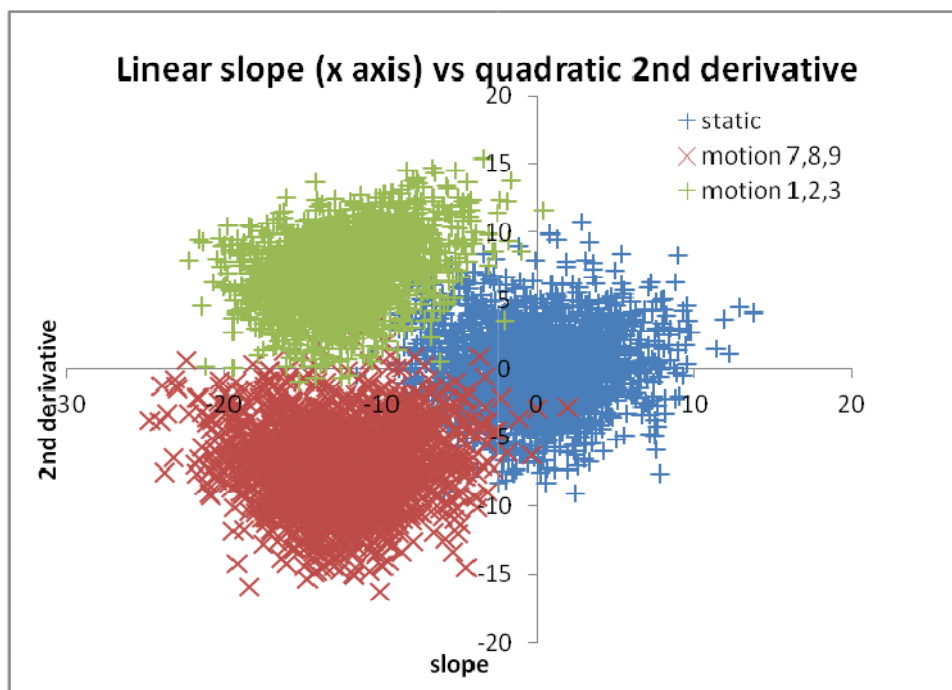


Figure 7. Linear regression slope vs. second derivative. Conditions same as in Figure 6. Red indicate motion just being seen in buffer. Green points indicate motion waning.

Using both the 2nd derivative and linear regression slope acts as a linear discriminator function. The data turns out to be slightly better at statistically indicating motion is occurring.

So What went Wrong?

As mentioned in the beginning, the project was considered a failure. Our house is L-shaped with the garage one leg of the "L". The driveway to the house is relatively narrow, such that a car must enter the garage at some angle non normal to the back of the garage. The front grill and bumper will not be normal to the centerline of the sensor (parallel to the back wall) until the car is most of the way into the garage. The same with slight less of an issue is true of the side wall.

The angle of detection of the HC-SR04 sensor is often quoted as 15-20 degrees. However, that angle depends on distance from the sensor, and the angle of the target face with respect to the sensor. See <http://www.fadstoobsessions.com/Learning-Electronics/Component-Testing/HC-SR04-Ultrasonic-Ranging.php>. For a more comprehensive report, see my repository <https://github.com/GaryDyr/HC-SR04-beam-tests>. Below is the angle data for a target with the face perpendicular to the imaginary centerline of the HC-SR04 sensor:




Figure 8. Detection angle with distance for a 2' x 4' plastic faced foam board, perpendicular to the centerline of the sensor.

The important point of this angle versus distance relationship is just how narrow the detection angle is with distance. If the size of the target is not very large, and not very close to perpendicular to the centerline of the sensor, there will be no detection.

At the start of the project, I had no portable logging system, i.e., a laptop with a reasonable microprocessor speed, to do real time logging of distance measurements on a car entering the garage. Only near the end of the project, a nearly ancient, but moderate cpu speed laptop, became available to use as a mobile logging station with my Arduino Uno.

The sensor was placed in one of two positions, relative to the Honda Pilot. Like most modern vehicles, it has a curved front and complex grill, with some portion of the bumper flat, but curving off to the sides. The first measurements were done near the "H" emblem height, which would have mostly impacted the web like grill. The second set of measurements were done with the sensor coincident with the license plate and bumper. The goal was two measurements, distance to the sensor and the velocity as a function of distance. The Arduino code for this was similar to that already discussed. The only changes were increase the maximum distance to 550 cm (18 ft). However, a completely new python module had to be written for the logging aspect through the USB serial port. See my repository <https://github.com/GaryDyr/Logging-Arduino-Data-HC-SR04>. Below is the distance profile from 9 runs pulling into the garage, with the start points matched up as best as possible:

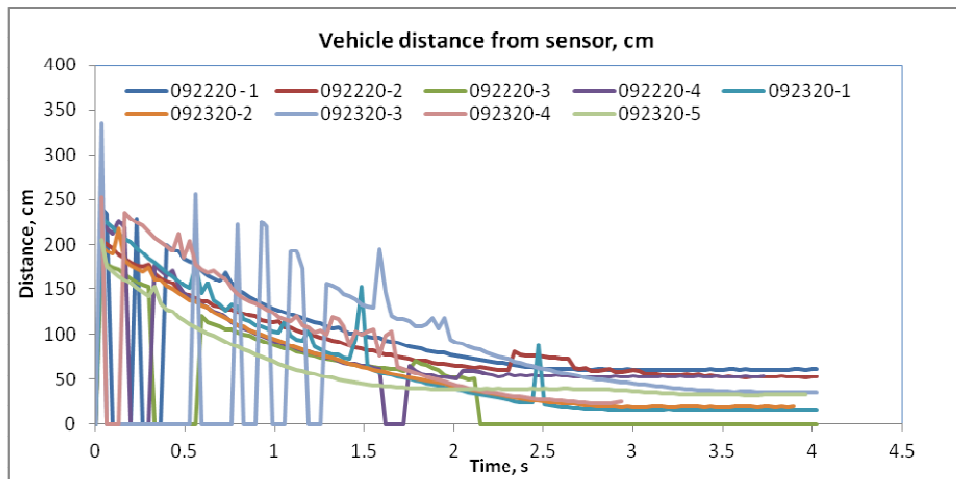


Figure 9. Car pulling into garage. Distances measured by sensor. Each curve time start has been adjusted to time car first registered a distance. Legend shows the date and number of the run. The 92320 data runs were with the sensor at license plate and bumper height.

There is a lot going on here. First, note that all the curves have been aligned to the first point that registered a measure of the car's distance from the sensor. At early times, when the car is far from the sensor, there are a large numbers of "dropout" values; the sensor found no return echo. This was likely due to multiple factors: 1. The angle of the car grill with respect to the microphone sensor was not within the microphone detection angle, so sonar pulses were reflected away from the microphone. 2. When occasional values do show up, they are probably stray reflections off some grill part that end up normal to the microphone. As the distance to the sensor closed, the car straightened out entering the garage with the car front becoming closer to perpendicular to the centerline of the sensor, and therefore the difference in reflected angle now fell within the detection angle of the microphone, e.g., Fig. 8.

This angle mismatch problem is also reflected in the distance measurements, echoes only start to show up at all, only 230 cm (7.5 ft) from sensor, which is only 42% of the expected 550 maximum distance range. That amounts to only 4 seconds worth of time, compared to the estimated 11 s to enter and park.

As would be expected, the velocity of the car is clearly decelerating rapidly over this short distance. Lower velocity changes would be expected beyond the 2.3 m we can see.

Velocities were really a problem to measure against the many dropouts. Velocity is measured as the distance difference between successive 33 ms ping interval. To get even a rough idea of the velocity/acceleration changes, Interpolation was necessary to get around the massive dropouts. The data was then averaged across all the trials:




Figure 10. Velocity change of Car pulling into garage, last 2.5 meters.

At the longest distances, the velocity is on the order of 134 cm/s, and dropping rapidly. So maximum velocity on entry into the garage could be 0-50% higher 149-200 cm/s. If echoes could have been received from longer distances, a sigmoid curve would be expected. Also, a calculation of the deceleration indicated it was not constant in the range we measured. The good news is that if we look back at the velocity range that a 5 element or 10 element buffer accommodates, motion detection will be possible down to very close to the stop distance limits for both buffers, but of course, with an increase in lag time. The 10 element buffer, would provide for even slower entry into the garage.

Further experiments tried to move the sensor so the beam was more normal to the bumper, as the car entered the garage at the typical angle it always did. No echo response could be raised, even though a 2'x3' plywood panel returned echoes at over 4.5 meters. As anticipated, with the plywood target, it was necessary to very carefully 'aim' the board at the microphone of the HC-SR04. The problem, therefore, is not the sensor. The shape of the target, i.e., the bumper, is too curved to send an echo with enough sound energy back to the microphone, especially with the narrow angle that the sensor reads at long distances.

Nevertheless, the initial project, developing a two sensor LED position monitoring system for car parking fails. It is not reasonable to adjust the position of a car into an area with only around 2 meters distance to maneuver. Can it ever work? Certainly, if a car is pulling straight into a garage, there is a much higher chance of getting the long range measurements needed. However, the shape of the front of the car is definitely an important factor. The more curved, the less likely to pick up long range reflections.

Sensor Code

Before the project died an ugly death at the hands of real garage sensor tests, I had already been doing some small scale tests with manually moving objects. The sketch code is not revolutionary, but some readers may find it useful as a starting point for other projects with multiple sensors, and maybe save someone else some time. Particularly, note that the side sensor coding sections have been commented out, so that only the front sensor readings were being used. The sketch is in the repository as: *usonic_progLED.ino*.

Other possible solutions

So what other possible solutions could be used? Of course, speculation is always fun and cheap. Here are few I came up with, but there are likely more.

The first was to abandon the idea of using a forward sensor, and just use the side sensor as a partial solution. Because the car side distance to sensor is smaller, though there still is an angle issue at entry start, as suggested by Figure 8 the angle of detection should improve. In my case, that likely would work fine for one car. The problem is that my garage has more room on one side that puts the detection distance at 1.5 m from the side of the car. At that distance, the detection angle is roughly 8 degrees from the centerline, so the prospect of early detection was iffy.

A series of PIR sensors, with adjustable range, or HC-SR04 sensors placed on the side wall might work. This requires at least 3 sensors placed along side walls at the distances for triggering different led color regions. For two cars this is 6 sensors. In addition, we are into many feet of wiring for remote sensors. Pushing this idea even further: To minimize wiring, ESP-01's or ESP32's connected with ultrasonic sensors could be use, but the combination of ESP and HC-SR04 or PIR sensors, puts heavy power demands on a battery pack. We need either many packs, or a single line power source. to power all the modules, but we run into all sorts of new issues with power management.

Another approach is to use combinations of different sensors. For example, some cheap Doppler radar sensors (e.g., HB100 or CDM324) can output the Doppler frequency, which can be converted to object velocity. If it is determined through many tests that the car velocity versus distance in the garage is consistent, then velocity could serve as a proxy for direct forward distance measurement. This could be coupled with an ultrasonic sensor to detect lateral vehicle distance. There are several big issues in using a Doppler sensor (for me at least). At this time, the sensors require a separate op amplifier circuit be constructed. There are some ready-made modules available, but they are relatively expensive. There are two stage or more op amp off the shelf modules that do work or might work, but there appear to be issues with filtering noise, op amp efficiency, and amplification factor, that have to be carefully negotiated. Also, off centerline measurements of velocity add an angle factor. At this time, getting away cheap with Doppler sensors requires sophisticated circuit knowledge and an oscilloscope to understand issues.